



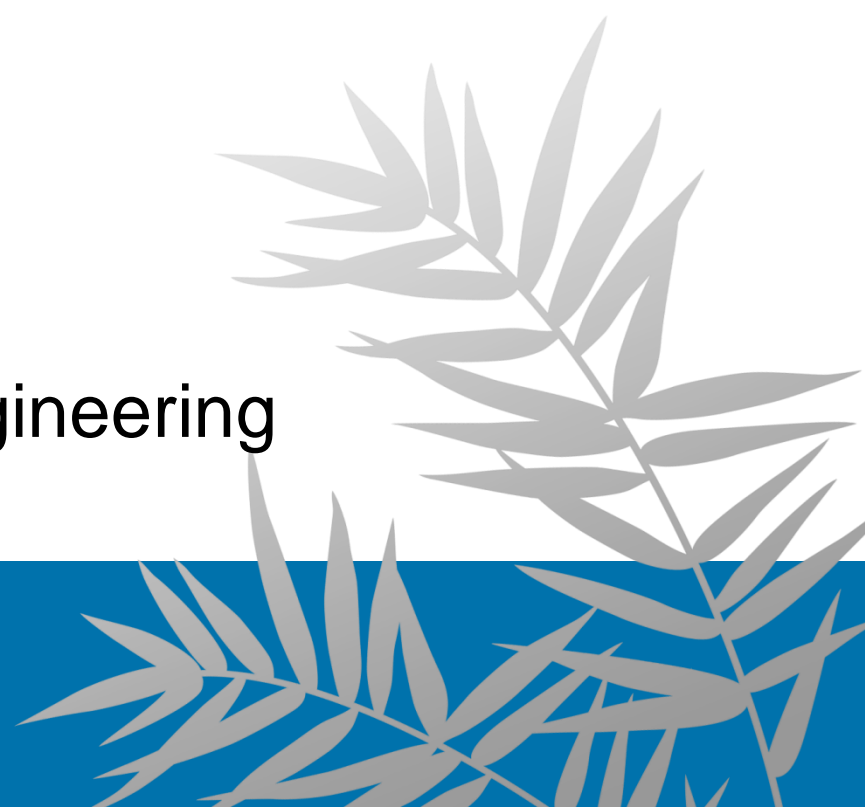
國立臺灣大學  
National Taiwan University

# CHAPTER 4

## GREEDY ALGORITHMS

Iris Hui-Ru Jiang  
Fall 2017

Department of Electrical Engineering  
National Taiwan University



# Outline

---

- Content:
  - Interval scheduling: The greedy algorithm stays ahead
  - Scheduling to minimize lateness: An exchange argument
  - Shortest paths
  - The minimum spanning tree problem
  - Implementing Kruskal's algorithm: Union-find
- Reading:
  - Chapter 4

# Greedy Algorithms

- An algorithm is **greedy** if it builds up a solution in small steps, **choosing a decision at each step myopically** to optimize some underlying criterion.
- It's **easy** to invent greedy algorithms for almost **any** problem.
  - Intuitive and fast
  - Usually not optimal
- It's **challenging** to prove greedy algorithms succeed in solving a nontrivial problem **optimally**.
  1. The greedy algorithm **stays ahead**.
  2. An **exchange** argument.


# Interval Scheduling

*The greedy algorithm stays ahead*



# The Interval Scheduling Problem

- Given: Set of requests  $\{1, 2, \dots, n\}$ ,  $i^{\text{th}}$  request corresponds to an interval with start time  $s(i)$  and finish time  $f(i)$ 
  - interval  $i$ :  $[s(i), f(i))$  ← requests don't overlap
- Goal: Find a compatible subset of requests of maximum size ← optimal



Maximum compatible subset  $\{2, 5, 8\}$

# Greedy Rule

- Repeat
  - Use a simple rule to select a first request  $i_1$
  - Once  $i_1$  is selected, reject all requests incompatible with  $i_1$ .
- Until run out of requests
- Q: How to decide a greedy rule for a good algorithm?
- A:


1. Earliest start time:  $\min s(i)$



2. Shortest interval:  $\min \{f(i) - s(i)\}$



3. Fewest conflicts:  $\min_{i=1..n} |\{j: j \text{ is not compatible with } i\}|$



4. Earliest finish time:  $\min f(i)$

# The Greedy Algorithm

- The 4<sup>th</sup> greedy rule leads to the optimal solution.

- We first accept the request that finish first
- Natural idea: Free resource ASAP

- The greedy algorithm:

$\emptyset$ : empty set =  $\{\}$

Interval-Scheduling( $R$ )

//  $R$ : undetermined requests;  $A$ : accepted requests

1.  $A = \emptyset$ ;
2. **while** ( $R$  is not empty) **do**
3.     **choose a request**  $i \in R$  **with minimum**  $f(i)$  **// greedy rule**
4.      $A = A + \{i\}$
5.      $R = R - \{i\} - X$ , where  $X = \{j: j \in R \text{ and } j \text{ is not compatible with } i\}$
6. **return**  $A$


- Q: Feasible?
- A: Yes! Line 5.

■  $A$  is a compatible set of requests.

- Q: Optimal? Efficient?

# The Greedy Algorithm Stays Ahead


- Q: How to prove optimality?
  - Let  $O$  be an optimal solution. Prove  $A = O$ ? or Prove  $|A| = |O|$ ?
- The greedy algorithm **stays ahead**.
  - We will compare the partial solutions of  $A$  and  $O$ , and show that the greedy algorithm is doing better in a step-by-step fashion.
- Let  $A$  be the output of the greedy algorithm,  $A = \{i_1, \dots, i_k\}$ , in the order they were added. Let  $O$  be the optimal solution,  $O = \{j_1, \dots, j_m\}$  in the ascending order of start (finish) times. For all indices  $r \leq k$ , we have  $f(i_r) \leq f(j_r)$ .
- Pf: Proof by induction!
  - Basis step: true for  $r = 1$ ,  $f(i_1) \leq f(j_1)$ .
  - Inductive step: hypothesis: true for  $r-1$ .
    - $f(i_{r-1}) \leq f(j_{r-1})$
    - $O$  is compatible,  $f(j_{r-1}) \leq s(j_r)$
    - Hence,  $f(i_{r-1}) \leq s(j_r)$ ;  $j_r \in R$  after  $i_{r-1}$  is selected in line 5.
    - According to line 3,  $f(i_r) \leq f(j_r)$ .





# The Greedy Algorithm Is Optimal

- The greedy algorithm returns an optimal set  $A$ .
- Pf: Proof by contradiction.
  - If  $A$  is not optimal, then an optimal set  $O$  must have more requests, i.e.,  $|O| = m > k = |A|$ .
  - Since  $f(i_k) \leq f(j_k)$  and  $m > k$ , there is a request  $j_{k+1}$  in  $O$ .
  - $f(j_k) \leq s(j_{k+1})$ ;  $f(i_k) \leq s(j_{k+1})$ .
  - Hence,  $j_{k+1}$  is compatible with  $i_k$ .  $R$  should contain  $j_{k+1}$ .
  - However, the greedy algorithm stops with request  $i_k$ , and it is only supposed to stop when  $R$  is empty.  $\rightarrow \leftarrow$



# Implementation: The Greedy Algorithm

Interval-Scheduling( $R$ )

//  $R$ : undetermined requests;  $A$ : accepted requests

1.  $A = \emptyset$
2. **while** ( $R$  is not empty) **do**
3.     **choose a request**  $i \in R$  with minimum  $f(i)$  // greedy rule
4.      $A = A + \{i\}$
5.      $R = R - \{i\} - X$ , where  $X = \{j: j \in R \text{ and } j \text{ is not compatible with } i\}$
6. **return**  $A$

- Running time: From  $O(n^2)$  to  $O(n \log n)$


- Initialization:

- $O(n \log n)$ : sort  $R$  in ascending order of  $f(i)$
    - $O(n)$ : construct  $S$ ,  $S[i] = s(i)$

- Lines 3 and 5:

- $O(n)$ : scan  $R$  once
    - We always select the first interval in  $R$
    - We do not delete all incompatible requests in line 5; we skip only those listed before the next selected interval.

# The Interval Scheduling Problem



Maximum compatible subset {1, 3, 5, 8}


# Interval Partitioning

*Interval coloring*




# What if We Have Multiple Resources?

- The **interval partitioning** problem:
  - a.k.a. the **interval coloring** problem: one resource = one color
  - Use as few resources as possible
- Given: Set of requests  $\{1, 2, \dots, n\}$ ,  $i^{\text{th}}$  request corresponds an interval with start time  $s(i)$  and finish time  $f(i)$ 
  - interval  $i$ :  $[s(i), f(i))$
- Goal: Partition these requests into a **minimum** number of **compatible** subsets, each corresponds to one resource



# How Many Resources Are Required?

- The **depth** of a set of intervals is the **maximum** number that pass over any single point on the time-line.
- In any instance of interval partitioning, the number of resources needed is **at least** the **depth** of the set of intervals.
- Pf:
  - Suppose a set of intervals has depth  $d$ , and let  $I_1, \dots, I_d$  all pass over a common point on the time-line.
  - Then each of these intervals must be scheduled on a different resource, so the whole instance needs **at least  $d$  resources**.



# Can We Reach the Lower Bound?

- The depth  $d$  is the lower bound on the number of required resources.
- Q: Can we always use  $d$  resources to schedule all requests?
- A: Yes.
- Q: How to prove the optimality?
- A:
  1. Find a bound that every possible solution must have at least a certain value
  2. Show that the algorithm under consideration always achieves this bound

# The Greedy Algorithm

- Assign a label to each interval. Possible labels:  $\{1, 2, \dots, d\}$ .
- Assign different labels for overlapping intervals.


## Interval-Partitioning( $R$ )

1.  $\{I_1, \dots, I_n\}$  = sort intervals in ascending order of their start times
2. **for**  $j$  **from** 1 **to**  $n$  **do**
3.     exclude the labels of all assigned intervals that are not compatible with  $I_j$
4.     **if** (there is a nonexcluded label from  $\{1, 2, \dots, d\}$ ) **then**
5.         assign a nonexcluded label to  $I_j$
6.     **else** leave  $I_j$  unlabeled

- Implementation:
  - Lines 3--5: find a resource compatible with  $I_j$ , assign this label
    - Record the finish time of the last added interval for each label
    - Compatibility checking:  $s(I_j) \geq f(\text{label}_j)$ 
      - Use priority queue to maintain labels



# The Interval Partitioning Problem



# Optimality (1/2)

- The greedy algorithm assigns every interval a label, and no two overlapping intervals receive the same label.
- Pf:
  1. No interval ends up unlabeled.
    - Suppose interval  $I_j$  overlaps  $t$  intervals earlier in the sorted list.
    - These  $t + 1$  intervals pass over a common point, namely  $s(I_j)$ .
    - Hence,  $t + 1 \leq d$ . Thus,  $t \leq d - 1$ ; at least one of the  $d$  labels is not excluded, and so there is a label that can be assigned to  $I_j$ .
      - i.e., line 6 never occurs!

## Interval-Partitioning( $R$ )

1.  $\{I_1, \dots, I_n\}$  = sort intervals in ascending order of their start times
2. **for**  $j$  **from** 1 **to**  $n$  **do**
3.     exclude the labels of all assigned intervals that are not compatible with  $I_j$
4.     **if** (there is a nonexcluded label from  $\{1, 2, \dots, d\}$ ) **then**
5.         assign a nonexcluded label to  $I_j$
6.     **else** leave  $I_j$  unlabeled

# Optimality (2/2)

- Pf: (cont'd)
  2. No two overlapping intervals are assigned with the same label.
    - Consider any two intervals  $I_i$  and  $I_j$  that overlap,  $i < j$ .
    - When  $I_j$  is considered (in line 2),  $I_i$  is in the set of intervals whose labels are excluded (in line 3).
    - Hence, the algorithm will not assign the label used for  $I_i$  to  $I_j$ .
  - Since the algorithm uses  $d$  labels, we can conclude that the greedy algorithm **always** uses the **minimum** possible number of labels, i.e., it is **optimal**!

## Interval-Partitioning( $R$ )

1.  $\{I_1, \dots, I_n\}$  = sort intervals in ascending order of their start times
2. **for**  $j$  **from** 1 **to**  $n$  **do**
3.     **exclude the labels of all assigned intervals that are not compatible with**  $I_j$
4.     **if** (there is a nonexcluded label from  $\{1, 2, \dots, d\}$ ) **then**
5.         assign a nonexcluded label to  $I_j$
6.     **else** leave  $I_j$  unlabeled

# Scheduling to Minimize Lateness

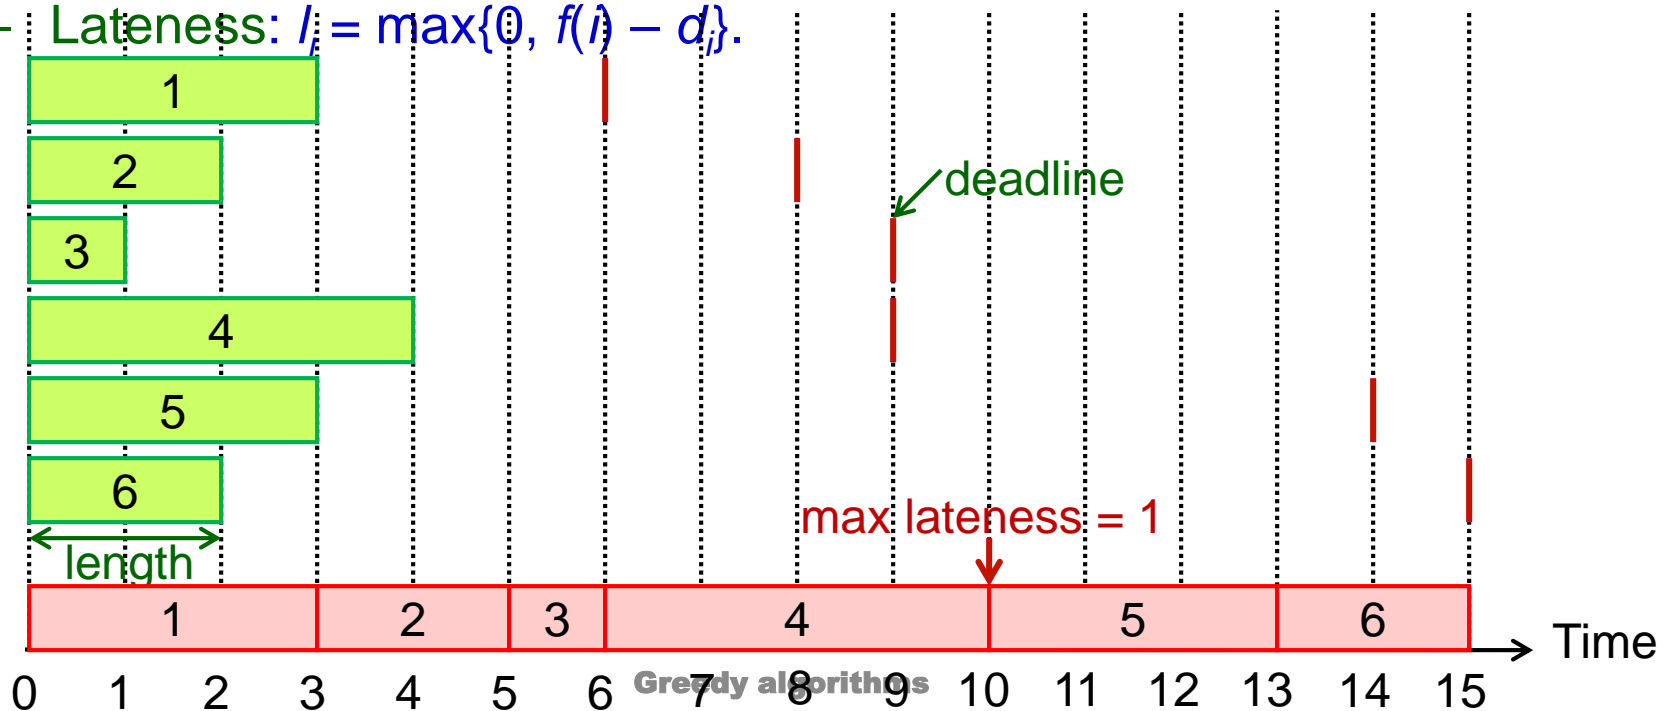
*An exchange argument*



# What If Each Request Has a Deadline?

- Given: A single resource is available starting at time  $s$ . A set of requests  $\{1, 2, \dots, n\}$ , request  $i$  requires a contiguous interval of length  $t_i$  and has a **deadline**  $d_i$ .
- Goal: Schedule all requests without overlapping so as to minimize the **maximum lateness**.

– Lateness:  $l_i = \max\{0, f(i) - d_i\}$ .



# Greedy Rule

- Consider requests in some order.
  - Shortest interval first: Process requests in ascending order of  $t_i$

	1	2
$t_i$	1	10
$d_i$	100	10

- Smallest slack: Process requests in ascending order of  $d_i - t_i$

	1	2
$t_i$	1	10
$d_i$	2	10

- Earliest deadline first: Process requests in ascending order of  $d_i$

Jobs with earlier deadlines  
get completed earlier


# Minimizing Lateness

- Greedy rule: Earliest deadline first!

Min-Lateness( $R, s$ )

//  $f$ : the finishing time of the last scheduled request

1.  $\{d_1, \dots, d_n\}$  = sort requests in ascending order of their deadlines
2.  $f = s$
3. **for**  $i$  **from** 1 **to**  $n$  **do**
4.     assign request  $i$  to the time interval from  $s(i) = f$  to  $f(i) = f + t_i$
5.      $f = f + t_i$
6. **return** the set of scheduled intervals  $[s(i), f(i))$  for all  $i = 1..n$



# No Idle Time

- Observation: The greedy schedule has no idle time.


– Line 4!

Min-Lateness( $R, s$ )

//  $f$ : the finishing time of the last scheduled request

1.  $\{d_1, \dots, d_n\}$  = sort requests in **ascending** order of their deadlines
2.  $f = s$
3. **for**  $i$  **from** 1 **to**  $n$  **do**
4.     assign request  $i$  to the time interval from  $s(i) = f$  to  $f(i) = f + t_i$
5.      $f = f + t_i$
6. **return** the set of scheduled intervals  $[s(i), f(i))$  for all  $i = 1..n$

- There is an optimal schedule with no idle time.





# No Inversions


- **Exchange argument:** Gradually transform an optimal solution to the one found by the greedy algorithm without hurting its quality.
- An **inversion** in schedule  $S$  is a pair of requests  $i$  and  $j$  such that  $s(i) < s(j)$  but  $d_j < d_i$ .
- All schedules without inversions and without idle time have the same maximum lateness.
- Pf:
  - If two different schedules have neither inversions nor idle time, then they can only differ in the order in which requests with identical deadlines are scheduled.
  - Consider such a deadline  $d$ . In both schedules, the jobs with deadline  $d$  are all scheduled consecutively (after all jobs with earlier deadlines and before all jobs with later deadlines).
  - Among them, the last one has the greatest lateness, and this lateness does not depend on the order of the requests.

# Optimality

- There is an optimal schedule with no inversions and no idle time.

- Pf:


- There is an optimal schedule  $O$  without idle time. (done!)
  1. If  $O$  has an inversion, there is a pair of jobs  $i$  and  $j$  such that  $j$  is scheduled immediately after  $i$  and has  $d_j < d_i$ .



2. After swapping  $i$  and  $j$  we get a schedule with one less inversion.
3. The new swapped schedule has a maximum lateness no larger than that of  $O$ .

$d_j$   
 $d_i$

Other requests have the same lateness



# Optimality: Exchange Argument

- Theorem: The greedy schedule  $S$  is optimal.
- Pf: Proof by contradiction
  - Let  $O$  be an optimal schedule with inversions.
  - Assume  $O$  has no idle time.
  - If  $O$  has no inversions, then  $S = O$ . done!
  - If  $O$  has an inversion, let  $i$ - $j$  be an adjacent inversion.
    - Swapping  $i$  and  $j$  does not increase the maximum lateness and strictly decreases the number of inversions.
    - This contradicts definition of  $O$ .

# Summary: Greedy Analysis Strategies

- An algorithm is **greedy** if it builds up a solution in small steps, **choosing a decision at each step myopically** to optimize some underlying criterion.
- It's **challenging** to prove greedy algorithms succeed in solving a nontrivial problem **optimally**.
  1. The greedy algorithm **stays ahead**: Show that after each step of the greedy algorithm, its partial solution is better than the optimal.
  2. An **exchange** argument: Gradually transform an optimal solution to the one found by the greedy algorithm without hurting its quality.

# Shortest Paths


*Edsger W. Dijkstra 1959*



# Edsger W. Dijkstra (1930—2002)

- 1972 Recipient of the ACM Turing Award

*The question of whether computers can think is as relevant as the question of whether submarines can swim.*




*If you want more effective programmers,  
you will discover that they should not waste their time debugging,  
they should not introduce the bugs to start with.*

*Program testing can be a very effective way to show the presence of bugs,  
but it is hopelessly inadequate for showing their absence.*


*-- Turing Award Lecture 1972, the humble programmer*

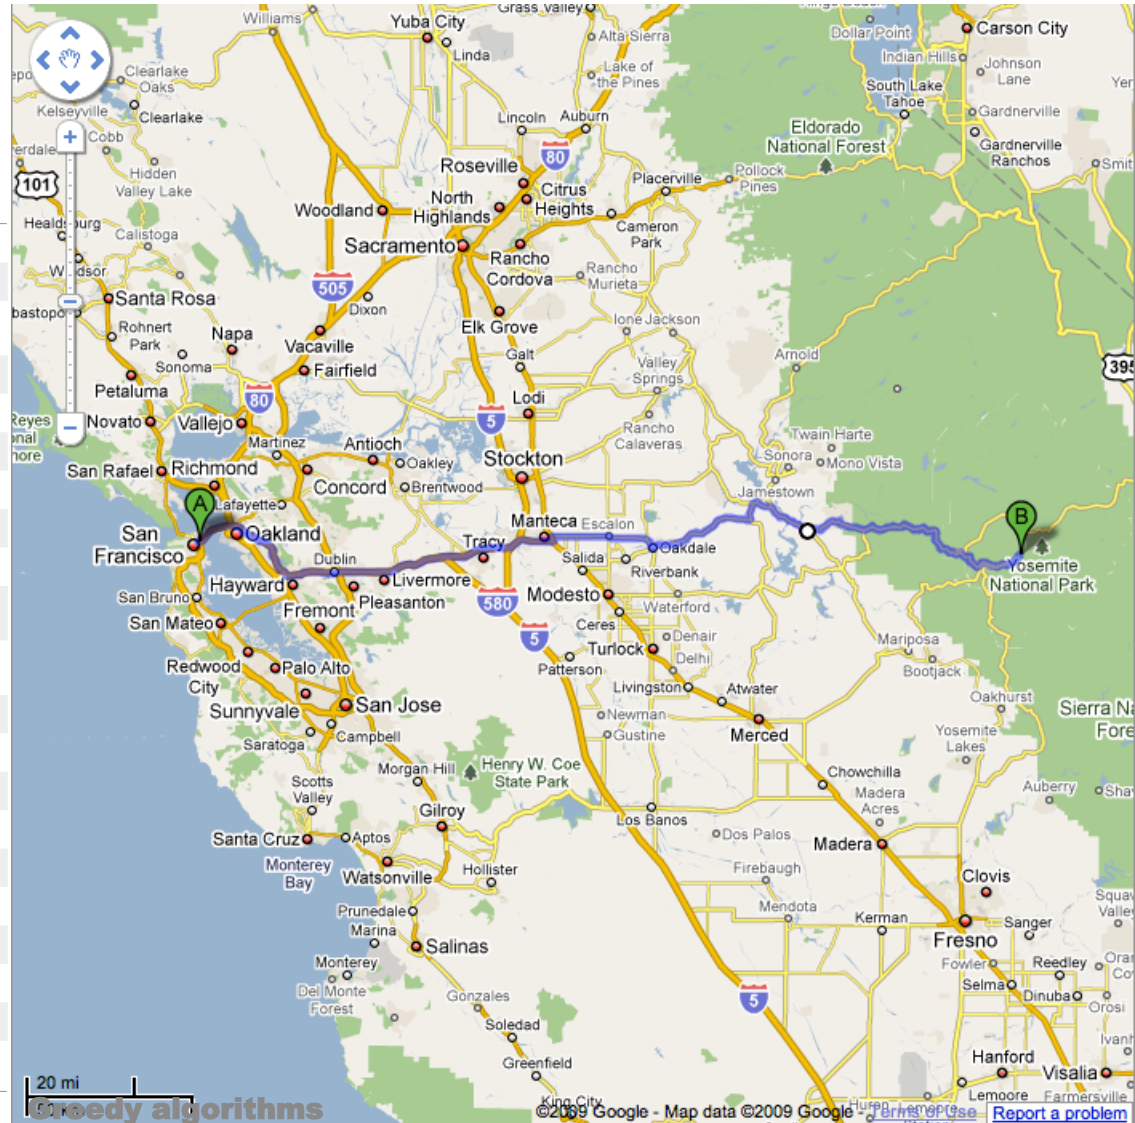
# Google Map

## Shortest path from San Francisco Shopping Centre to Yosemite National Park

 Westfield San Francisco Center  
845 Market St, San Francisco, CA 94103-1923 - (415) 495-5656

1. Head **northeast** on **Mission St** toward **4th St**  
About 2 mins  
go 0.6 mi  
total 0.6 mi
2. Turn **right** at **1st St**  
About 2 mins  
go 0.4 mi  
total 1.0 mi
3. Slight **right** to merge onto **Bay Bridge/I-80 E** toward **Oakland**  
Continue to follow I-80 E  
About 6 mins  
go 5.9 mi  
total 6.8 mi
4. Take the exit onto **I-580 E** toward **CA-24/Hayward/Stockton**  
About 42 mins  
go 46.2 mi  
total 53.0 mi
5. Continue onto **I-205 E**  
About 14 mins  
go 14.6 mi  
total 67.6 mi
6. Merge onto **I-5 N**  
About 1 min  
go 0.8 mi  
total 68.4 mi
7. Take exit **461** to merge onto **CA-120 E** toward **Manteca/Sonora**  
About 6 mins  
go 6.4 mi  
total 74.8 mi
8. Take the exit onto **CA-120 E/CA-99 N** toward **Sacramento/Sonora N**  
About 2 mins  
go 1.7 mi  
total 76.5 mi
9. Take exit **242** for **CA-120 E/Yosemite Ave** toward **Sonora**  
go 0.2 mi  
total 76.7 mi
10. Turn **right** at **CA-120 E/Yosemite Ave**  
Continue to follow CA-120 E  
About 33 mins  
go 19.8 mi  
total 96.5 mi
11. Turn **left** at **CA-108 E/CA-120 E/E F St**  
Continue to follow CA-108 E/CA-120 E  
About 38 mins  
go 25.0 mi  
total 121 mi
12. Turn **right** at **CA-120 E**  
About 20 mins  
go 12.6 mi  
total 134 mi
13. Turn **right** at **Old Priest Grde**  
About 5 mins  
go 1.8 mi  
total 136 mi
14. Slight **left** at **CA-120 E**  
About 53 mins  
go 34.8 mi  
total 171 mi
15. Slight **right** at **Big Oak Rd**  
About 17 mins  
go 9.4 mi  
total 180 mi
16. Turn **left** at **CA-140 E/EI Portal Rd** (signs for **Yosemite Valley/CA-41/Fresno**)  
About 2 mins  
go 0.9 mi  
total 181 mi
17. Take the **1st right** onto **CA-140 E**  
About 8 mins  
go 5.1 mi  
total 186 mi
18. Turn **left** at **Sentinel Dr**  
About 2 mins  
go 0.3 mi  
total 186 mi
19. Continue onto **Village Dr**  
About 1 min  
go 0.2 mi  
total 187 mi
20. **Village Dr** turns slightly **right** and becomes **Ahwahnee Dr**  
go 364 ft  
total 187 mi

 Yosemite National Park  
9039 Village Dr, Yosemite Natl Pk, CA 95389 - (209) 372-0200





# Floor Guide in a Shopping Mall

- Direct shoppers to their destinations in real-time

**Westfield San Francisco Centre**

Home | Product Search | Gift Cards | Fashion | Sign-Up & Win

Stores | Dining | Movies | Hours | Directions | Services | Offers & Events


Home > Find a Store > Abercrombie & Fitch

### Find a Store

Enter a store name

**All Stores**


- Books, Cards & Gifts
- Children & Teen's Fashions & Accessories
- Department Stores
- Electronics & Music
- Entertainment
- Fashion Accessories & Jewelry
- Fine Jewelry & Watches
- Health & Beauty
- Home Furnishings & Accessories
- Leathers, Luggage & Handbags
- Men's Fashion
- Services
- Shoes
- Sportswear & Equipment
- Sunglasses & Optical
- Wireless Communications & Services
- Women's Fashions



**Abercrombie & Fitch**  
Men's Fashion , Women's Fashions, All Stores  
Abercrombie and Fitch includes authentic vintage clothing for casual teens and young adults. This upscale teen apparel store has premium denim, as well as polos and t-shirts.

Level: 1 - Street  
Store #: 67

Phone: 415.284.9276  
Web: Visit Abercrombie & Fitch now



The map shows the layout of Level 1 (Street Level) of the Westfield San Francisco Centre. It includes major streets like Market Street, Mason Street, and Leavenworth Street. A large blue area is labeled 'bloomingdales'. A white arrow points to the location of Abercrombie & Fitch, which is marked with a red 'A' and the number 67. Other store locations are marked with red numbers and letters.



# The Shortest Path Problem

- Given:

- Directed graph  $G = (V, E)$

- Length  $l_e$  = length of edge  $e = (u, v) \in E$

- Distance; time; cost

- $l_e \geq 0$

- Q: what if undirected?


- A: 1 undirected edge = 2 directed ones

- Source  $s$

- Goal:

- Shortest path  $P_v$  from  $s$  to each other node  $v \in V - \{s\}$

- Length of path  $P$ :  $l(P) = \sum_{e \in P} l_e$



$$l(a \rightarrow b) = l(1 \rightarrow 3 \rightarrow 6 \rightarrow 5) \\ = 9 + 2 + 9 = 20$$

# Dijkstra's Algorithm

Dijkstra( $G, l$ )

//  $S$ : the set of **explored** nodes

// for each  $u \in S$ , we store a **shortest path distance**  $d(u)$  from  $s$  to  $u$


1. initialize  $S = \{s\}$ ,  $d(s) = 0$

2. **while**  $S \neq V$  **do**

3.     select a node  $v \notin S$  with at least one edge from  $S$  for which

4.          $d'(v) = \min_{e=(u,v): u \in S} (d(u) + l_e)$  ←

5.     add  $v$  to  $S$  and define  $d(v) = d'(v)$



shortest path to some  $u$  in explored part, followed by a single edge  $(u, v)$

$$d'(x) = 2; d'(y) = 4; d'(z) = 5$$

$$d''(y) = 3; d''(z) = 4$$

# Correctness

The algorithm  
stays ahead!


Q: Why  $l_e \geq 0$ ?

- Loop invariant: Consider the set  $S$  at any point in the algorithm's execution. For each node  $u \in S$ ,  $d(u)$  is the length of the shortest  $s$ - $u$  path  $P_u$ .

- Pf: **Proof by induction on  $|S|$**

- Basis step: trivial for  $|S| = 1$ .
- Inductive step: hypothesis: true for  $k \geq 1$ .

- Grow  $S$  by adding  $v$ ;  
let  $(u, v)$  be the final edge on our  $s$ - $v$  path  $P_v$ .
- By induction hypothesis,  $P_u$  is the shortest  $s$ - $u$  path.
- Consider any other  $s$ - $v$  path  $P$ ;  $P$  must leave  $S$  somewhere; let  $y$  be the first node on  $P$  that is not in  $S$ , and  $x \in S$  be the node just before  $y$ .
- $P$  cannot be shorter than  $P_v$  because it is already at least as long as  $P_v$  by the time it has left the set  $S$ .
- At iteration  $k+1$ ,  $d(v) = d'(v) = d(u) + l_{e=(u,v)} \leq d(x) + l_{e=(x,y)} \leq l(P)$




# Implementation


- Q: How to do line 4 efficiently?  
$$d'(v) = \min_{e = (u, v): u \in S} d(u) + l_e$$
- A: Explicitly maintain  $d'(v)$  in the view of each unexplored node  $v$  instead of  $S$ 
  - Next node to explore = node with minimum  $d'(v)$ .
  - When exploring  $v$ , update  $d'(w)$  for each outgoing  $(v, w)$ ,  $w \notin S$ .
- Q: How?

Operation	Dijkstra	Array	Binary heap	Fibonacci heap
Insert				
ExtractMin				
ChangeKey				
IsEmpty				
Total				


# Example




(a)




(b)




(c)



(d)



(e)



(f)

# Recap Heaps: Priority Queues

*Binary Tree Application*



# Priority Queue

- In a priority queue (PQ)
  - Each element has a priority (key)
  - Only the element with **highest** (or **lowest**) priority can be deleted
    - **Max** priority queue, or **min** priority queue
  - An element with arbitrary priority can be inserted into the queue at any time

Operation	Binary heap	Fibonacci heap
FindMin	$\Theta(1)$	$\Theta(1)$
ExtractMin	$\Theta(\lg n)$	$O(\lg n)$
Insert	$\Theta(\lg n)$	$\Theta(1)$
ChangeKey	$\Theta(\lg n)$	$\Theta(1)$

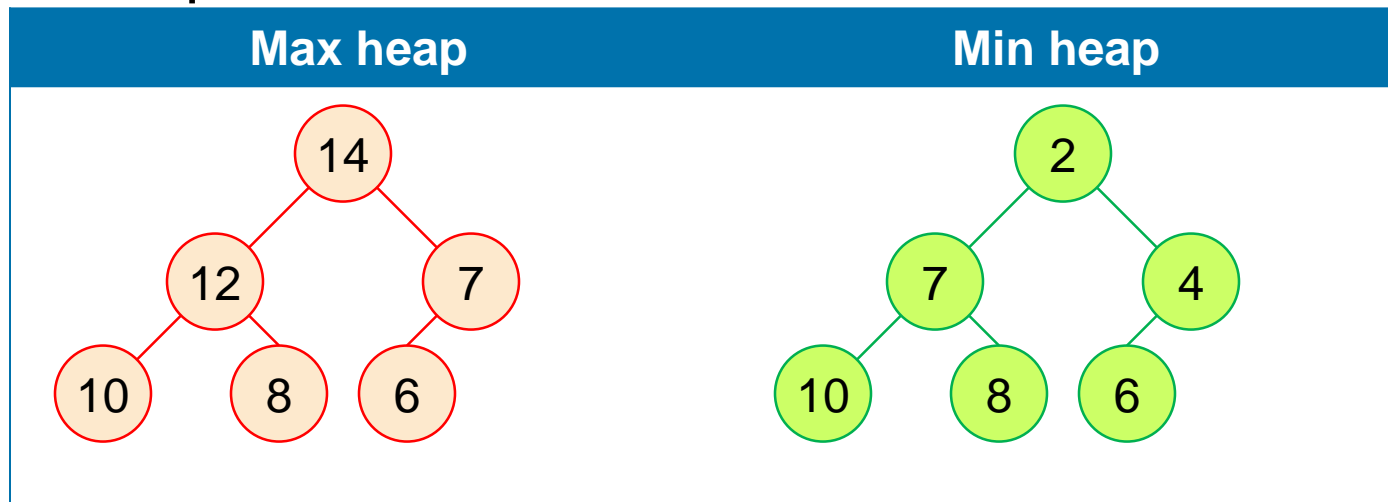
The time complexities are worst-case time for binary heap, and amortized time complexity for Fibonacci heap

Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein  
*Introduction to Algorithms, 2<sup>nd</sup> Edition*. MIT Press and McGraw-Hill, 2001.

Fredman M. L. & Tarjan R. E. (1987). Fibonacci heaps and their uses in improved network optimization algorithms. *Journal of the ACM* 34(3), pp. 596-615.

# Heap


- Definition: A **max (min) heap** is
  - A **max (min)** tree:  $key[parent] \geq (\leq) key[children]$
  - A complete binary tree
- Corollary: Who has the **largest** (smallest) key in a max (min) heap?
  - Root!
- Example





# Class *MaxHeap*

- Implementation?
  - Complete binary tree  $\Rightarrow$  array representation




	[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
heap	-	14	12	7	10	8	6	-	-	-	-



# Insertion into a Max Heap (1/3)

- Maintain heap property all the times


- *Push(1)*



[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
-	20	15	2	14	10	-	-	-	-	-



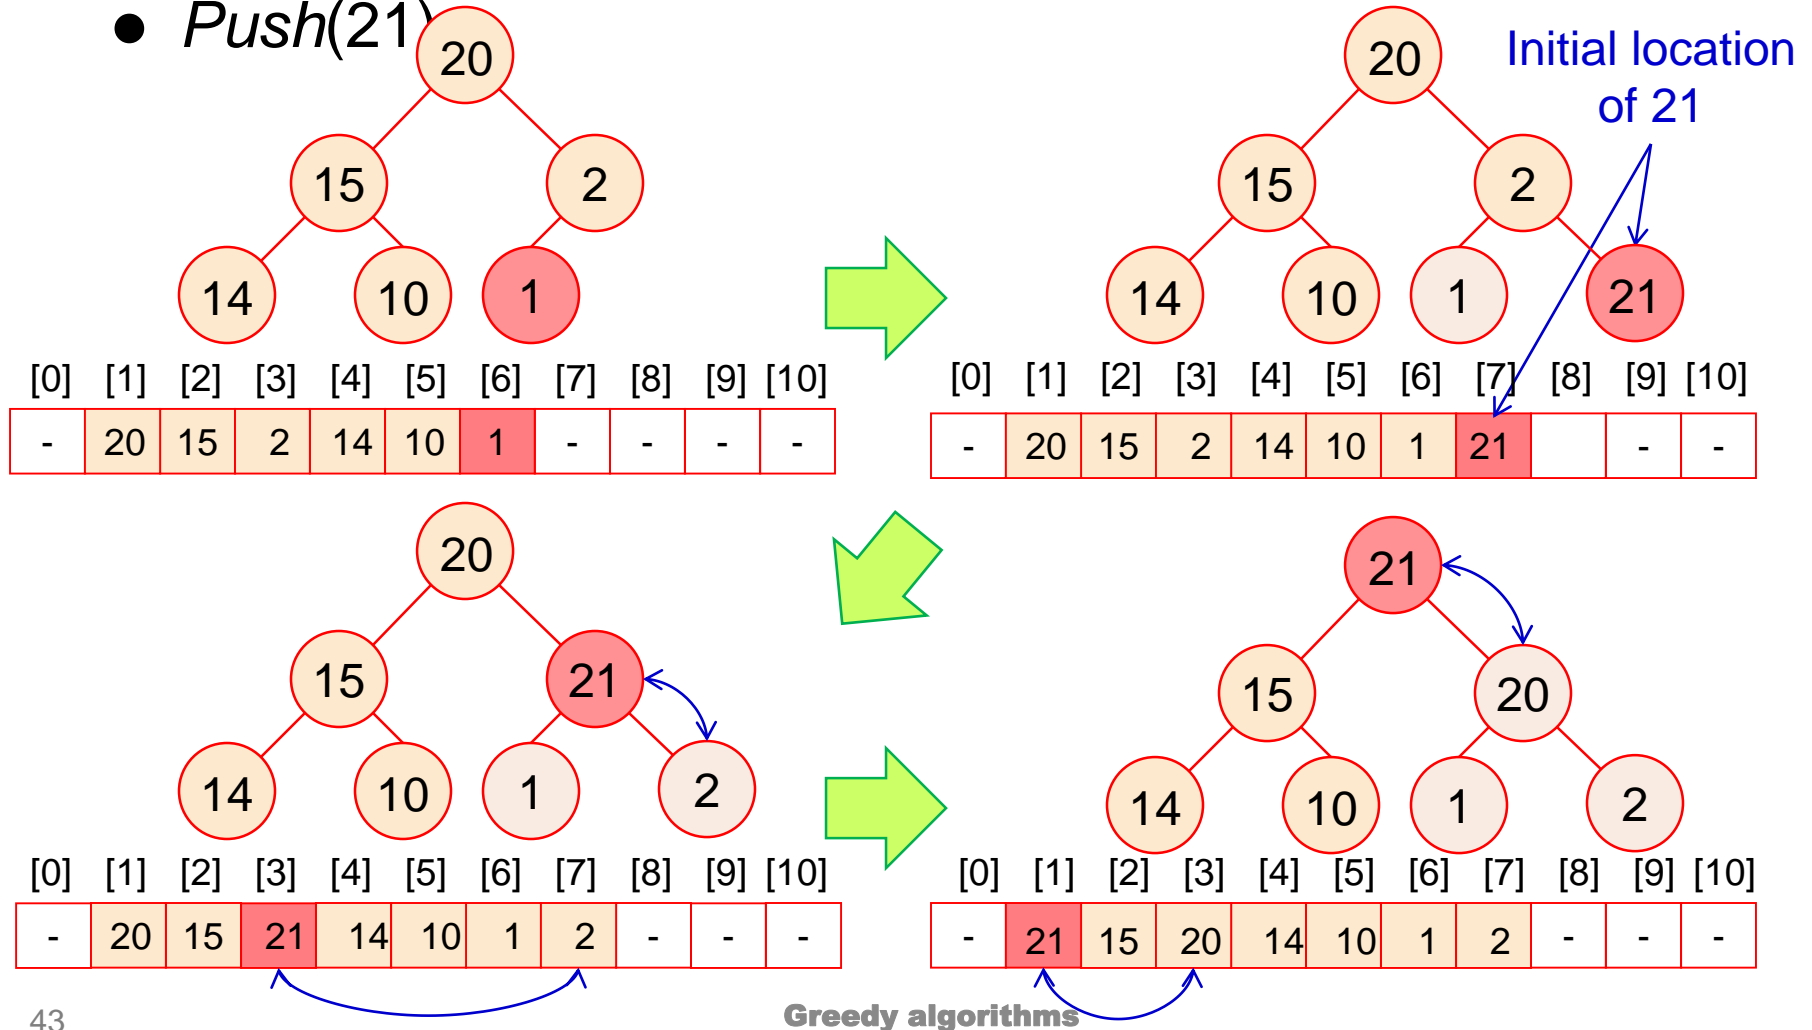
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
-	20	15	2	14	10		-	-	-	-



[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
-	20	15	2	14	10	1	-	-	-	-

# Insertion into a Max Heap (2/3)

- Maintain heap  $\Rightarrow$  bubble up if needed!
- *Push(21)*




# Insertion into a Max Heap (3/3)

- Time complexity?
  - How many times to bubble up in the worst case?
  - Tree height:  $\Theta(\lg n)$

# Deletion from a Max Heap (1/3)

- Maintain heap  $\Rightarrow$  trickle down if needed!


- *Pop()*



# Deletion from a Max Heap (2/3)

- Maintain heap  $\Rightarrow$  trickle down if needed!

- *Pop()*



# Deletion from a Max Heap (3/3)

---

- Time complexity?
  - How many times to trickle down in the worst case?  $\Theta(\lg n)$

# Max Heapify

- **Max** (**min**) heapify = maintain the **max** (**min**) heap property

- What we do to trickle down the root in deletion






- Assume  $i$ 's left & right subtrees are heaps

- But  $key[i]$  may be  $< (>)$   $key[children]$

- Heapify  $i$  = trickle down  $key[i]$

⇒ the tree rooted at  $i$  is a heap

$heapSize = 5$   
 $capacity = 10$

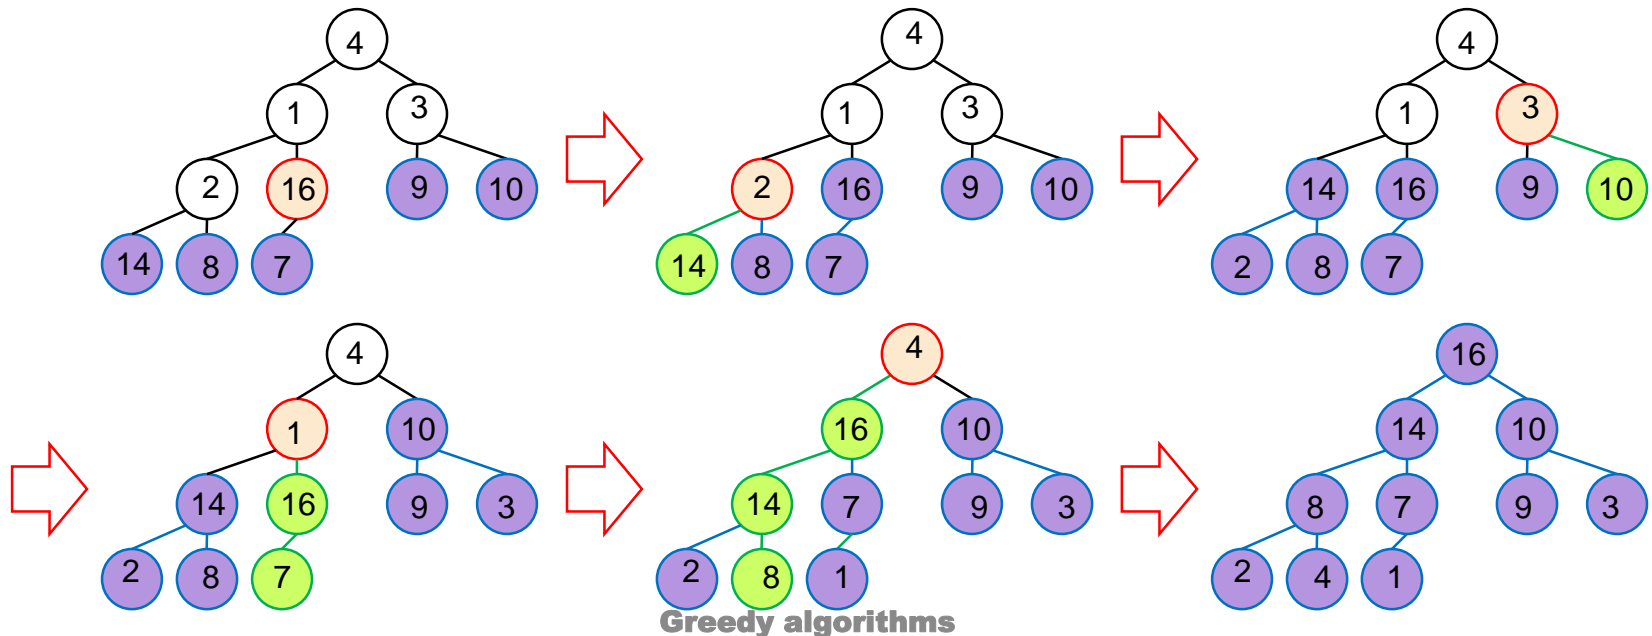




# How to Build a Max Heap?

- How to convert any complete binary tree to a max heap?
- Intuition: **Max heapify in a bottom-up manner**
  - Induction basis: Leaves are already heaps
  - Inductive steps: Start at parents of leaves, work upward till root
  - Time complexity:  $O(n \lg n)$

	[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
a	-	4	1	3	2	16	9	10	14	8	7



Greedy algorithms

# Minimum Spanning Trees

*Robert C. Prim 1957 (Dijkstra 1959)*


*Joseph B. Kruskal 1956*

*Reverse-delete*




# Minimum Spanning Graphs

- Q: How can a cable TV company lay cable to a new neighborhood, of course, as cheaply as possible?
- A: Curiously and fortunately, this problem is a case where many greedy algorithms optimally solve.
- Given
  - Undirected graph  $G = (V, E)$ 
    - Nonnegative cost  $c_e$  for each edge  $e = (u, v) \in E$ 
      - $c_e \geq 0$
- Goal
  - Find a subset of edges  $T \subseteq E$  so that
    - The subgraph  $(V, T)$  is connected
    - Total cost  $\sum_{e \in T} c_e$  is minimized




# Minimum Spanning ?????

- Q: Let  $T$  be a minimum-cost solution. What should  $(V, T)$  be?
- A:
  - By definition,  $(V, T)$  must be connected.
  - We show that it also contains no cycles.
  - Suppose it contained a cycle  $C$ , and let  $e$  be any edge on  $C$ .
  - We claim that  $(V, T - \{e\})$  is still connected
  - Any path previously used  $e$  can now go path  $C - \{e\}$  instead.
  - It follows that  $(V, T - \{e\})$  is also a valid solution, but cheaper.
  - Hence,  $(V, T)$  is a tree.
- Goal
  - Find a subset of edges  $T \subseteq E$  so that
    - $(V, T)$  is a tree,
    - Total cost  $\sum_{e \in T} c_e$  is minimized.




# Greedy Algorithms (1/3)




- Q: What will you do?
- All three **greedy** algorithms produce an MST.
- Kruskal's algorithm:
  - Start with  $T = \{\}$ .
  - Consider edges in ascending order of cost.
  - Insert edge  $e$  in  $T$  as long as it does not create a cycle; otherwise, discard  $e$  and continue.

# Greedy Algorithms (2/3)



- Q: What will you do?
- All three **greedy** algorithms produce an MST.
- Prim's algorithm: (c.f. Dijkstra's algorithm)
  - Start with a root node  $s$ .
  - Greedily grow a tree  $T$  from  $s$  outward.
  - At each step, add the cheapest edge  $e$  to the partial tree  $T$  that has exactly one endpoint in  $T$ .

# Greedy Algorithms (3/3)



- Q: What will you do?
- All three **greedy** algorithms produce an MST.
- Reverse-delete algorithm: (reverse of Kruskal's algorithm)
  - Start with  $T = E$ .
  - Consider edges in descending order of cost.
  - Delete edge  $e$  from  $T$  unless doing so would disconnect  $T$ .


# Cut Property (1/3)

- Simplifying assumption: All edge costs  $c_e$  are distinct.
- Q: When is it safe to include an edge in the MST?
- Cut Property: Let  $S$  be any subset of nodes, and let  $e = (v, w)$  be the minimum cost edge with one end in  $S$  and the other in  $V-S$ . Then every MST contains  $e$ .
- Pf: **Exchange argument!**
  - Let  $T$  be a spanning tree that does not contain  $e$ . We need to show that  $T$  does not have the minimum possible cost.
  - Since  $T$  is a spanning tree, it must contain an edge  $f$  with one end in  $S$  and the other in  $V-S$ .
  - Since  $e$  is the cheapest edge with this property, we have  $c_e < c_f$ .
  - Hence,  $T - \{f\} + \{e\}$  is a spanning tree that is cheaper than  $T$ .
- Q: What's wrong with this proof?
- A: Take care about the definition!



# Cut Property (2/3)

- Q: What's wrong with this proof?
- A: Take care about the definition!
  - Spanning tree: connected & acyclic



# Cut Property (3/3)

- Cut Property: Let  $S$  be any subset of nodes, and let  $e = (v, w)$  be the minimum cost edge with one end in  $S$  and the other in  $V-S$ . Then every MST contains  $e$ .

- Pf: **Exchange argument!**

- Let  $T$  be a spanning tree that does not contain  $e$ .

- $T$  is a spanning tree;  $\exists$  path  $P \in T$  from  $v$  to  $w$

- Let  $e' = (v', w')$  on  $P$ ,  $v' \in S$  and  $w' \in V-S$ .

- $T' = T - \{e'\} + \{e\}$  is a spanning tree


- $(V, T')$  must be connected:

$(V, T)$  is connected, any path in  $(V, T)$  using  $e'$  can be rerouted in  $(V, T')$  by  $v' \rightarrow v$ ,  $(v, w)$ ,  $w \rightarrow w'$ .

- $(V, T')$  must be acyclic:

The only cycle in  $(V, T' + \{e'\})$  is  $e + P$ , it isn't in  $(V, T)$

- Since  $c_e < c_{e'}$ ,  $T'$  is cheaper than  $T$ .



# Cycle Property

Optimality of Reverse-delete algorithm!

- Q: When is it safe to exclude an edge out?
- Cycle Property: Let  $C$  be any cycle in  $G$ , and let  $e = (v, w)$  be the maximum cost edge in  $C$ . Then  $e$  does not belong to any MST.
- Pf: **Exchange argument!** (Similar to Cut Property)

# Implementing MST Algorithms


*Priority queue*

*Union-find*



# Prim's Example

- R. C. Prim, 1957
- Procedure:
  - Start with a root node  $s$ .
  - Greedily grow a tree  $T$  from  $s$  outward.
  - At each step, add the cheapest edge  $e$  to the partial tree  $T$  that has exactly one endpoint in  $T$ .



# Prim's Algorithm

Dijkstra( $G, l$ )

//  $S$ : the set of **explored** nodes

// for each  $u \in S$ , we store a **shortest path distance**  $d(u)$  from  $s$  to  $u$

1. initialize  $S = \{s\}$ ,  $d(s) = 0$

2. **while**  $S \neq V$  **do**

3.     select a node  $v \notin S$  with at least one edge from  $S$  for which

4.          $d'(v) = \min_{e = (u, v): u \in S} d(u) + l_e$


5.     add  $v$  to  $S$  and define  $d(v) = d'(v)$

- Q: How to change Dijkstra's algorithm to Prim's?
- Q: How to implement?

R. C. Prim: *Shortest connection networks and some generalizations*.  
In *Bell System Technical Journal*, 36 (1957), pp. 1389–1401.

E. W. Dijkstra: *A note on two problems in connexion with graphs*.  
In *Numerische Mathematik*, 1 (1959), S. pp. 269–271.

# Kruskal's Algorithm



- J. B. Kruskal, 1956
- Procedure:
  - Start with  $T = \{\}$ .
  - Consider edges in ascending order of cost.
  - Insert edge  $e$  in  $T$  as long as it does not create a cycle; otherwise, discard  $e$  and continue.



Kruskal( $G, c$ )

1.  $\{e_1, e_2, \dots, e_m\}$  = sort edges in ascending order of their costs
2.  $T = \{\}$
3. **for each**  $e_i = (u, v)$  **do**
4.     **if** ( $u$  and  $v$  are not connected by edges in  $T$ ) **then** // different subtrees
5.          $T = T + \{e_i\}$  // merge these two corresponding subtrees

J. B. Kruskal: *On the shortest spanning subtree of a graph and the traveling salesman problem*. In *Proceedings of the American Mathematical Society*, 7(1) (Feb, 1956), pp. 48–50.

# The Union-Find Data Structure (1/2)


- Union-find data structure represents disjoint sets
  - Disjoint sets: elements are disjoint
  - Each set has a representative
  - Operations:
    - MakeUnionFind( $S$ ): initialize a set for each element in  $S$
    - Find( $u$ ): return the representative of the set containing  $u$
    - Union( $A, B$ ): merge sets  $A$  and  $B$





# The Union-Find Data Structure (2/2)

- Implementation: **disjoint-set forest**
  - Representative is the **root**; link: from children to parent
  - Union: attach the smaller to the larger one (**union by rank**)
  - Find: trace back to root and redirect the link (**path compression**)
- Running time: **union by rank + path compression**
  - The **amortized** running time per operation is  $O(\alpha(n))$ ,  $\alpha(n) < 5$  !!
    - Average running time of a sequence of  $n$  operations



B. A. Galler & M. J. Fischer. An improved equivalence algorithm.  
*Comm. of the ACM*, 7(5), (May 1964), pp. 301–303.

R. E. Tarjan & J. van Leeuwen. Worst-case analysis of set union algorithms.  
*Journal of the ACM*, 31(2), pp. 245–281, 1984.

# Implementing Kruskal's Algorithm

Kruskal( $G, c$ )

1.  $\{e_1, e_2, \dots, e_m\}$  = sort edges in ascending order of their costs
2.  $T = \{\}$
3. **for each**  $e_i = (u, v)$  **do**
4.     **if** ( $u$  and  $v$  are not connected by edges in  $T$ ) **then** // different subtrees
5.          $T = T + \{e_i\}$  // merge these two corresponding subtrees

- Use the union-find data structure
  - Maintain a disjoint set for each connected component (subtree)
  - Line 1: sort edge costs
  - Line 4: “Find” twice for each edge (total  $m$  edges in  $G$ )
  - Line 5: “Union” possibly once for each edge (total  $n-1$  edges in  $T$ )
  - Comparison sort + simple disjoint set:  $O(m \log n)$
  - Linear sort + union-find:  $O(m \alpha(m, n))$